Introduction to Algorithms and Data Structures

Bachmann–Landau notation

Giacomo Bergami

Homogeneous linear recurrence relations with constant coefficients

Objectives

- Understanding how to estimate the computational complexity for recursive algorithms without the **Master Theorem**.
- Understanding the difference between some Bachmann–Landau notations' operators (O and Θ).
- Understanding how memoization and caching could considerably decrease the computational complexity.
- Being able to generate recurrence relations from algorithms and vice-versa.

Bunny Island

Rabbit Island in Japan.

Leonardo da Pisa (also known as Fibonacci) was interested in many things, including the following population dynamics problem:

How quickly would a rabbit population grow under appropriate conditions?

In particular, starting from a pair of rabbits on a desert island, how many couples would you have in n years?

Outlining the problem

- A pair of rabbits generates two bunnies each year.
- Rabbits only start breeding in the second year after their birth.
- Rabbits are immortal.

Which is the resulting formula?

Fibonacci Sequence

In the n-th year, you still have all the previous year's couples, plus two new bunnies for each pair that was there two years ago.

$$\forall n \in \mathbb{N}\backslash\{0\}.\, F(n) = \begin{cases} 1 & n \leq 2 \\ F(n-1) + F(n-2) & otherwise \end{cases}$$

Let's Code!

See the implementation at `https://github.com/jackbergus/`
`NotesOnProgramming2020/tree/master/src/big_o_notation`.

Formal proof for the Fibonacci Sequence's correctness:

- $ai + bj + ck$ is shorthand for "there are a pairs of sexually active bunnies, b pairs of 1-year old bunnies, and c pairs of newborns".
- In the next year, we will have $(a + b)i + ck + (a + b)k$ bunnies.
- In the previous year, there were b pairs of new borns, there were b pairs of sexually active bunnies, so that $bi + (a - b)j + bk$.
- $(a + b) + a + (c + b) = (a + b + c) + (a + b)$

Bachmann-Landau notations (O and Θ, 1)

$$f(n) = \Theta(g(n)) \qquad f(n) = O(g(n))$$
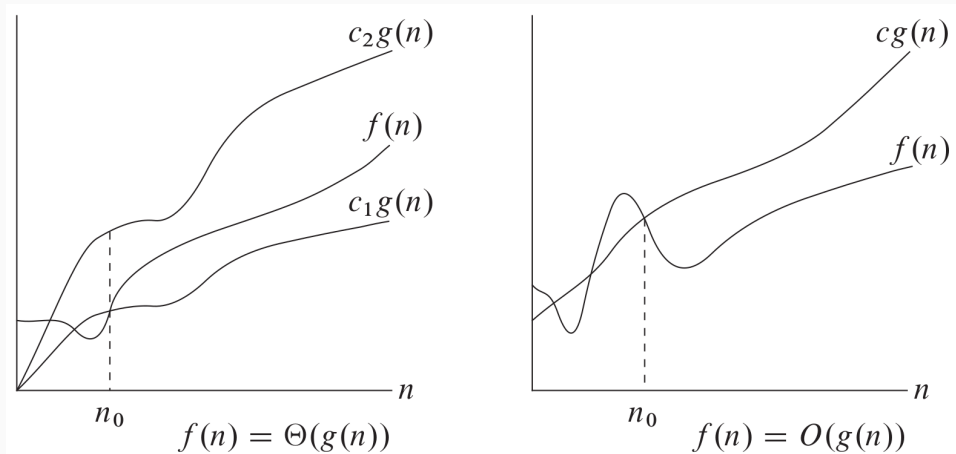
Bachmann–Landau notations (O and Θ, 2)

- If an algorithm F is of $\Theta(T(n))$, it means that the running time of the algorithm as n (input size) gets larger is proportional to $T(n)$.

$$0 < \lim_{n \to \infty} \frac{F(n)}{T(n)} < \infty \quad \exists k_1, k_2 > 0, n_0.\ \forall n > n_0.\ k_1 T(n) \leq F(n) \leq k_2 T(n)$$

- If an algorithm F is of $O(T(n))$, it means that the running time of the algorithm as n gets larger is at most proportional to $T(n)$.

$$\lim_{n \to \infty} \frac{F(n)}{T(n)} < \infty \qquad \exists k > 0, n_0.\ \forall n > n_0.\ |F(n)| \leq kT(n)$$

Observations on O and Θ

- If the algorithm has computational complexity
  $O(T_1(n) + T_2(n))$ where $\forall n.T_1(n) \geq T_2(n)$, then
  $T_1(n) + T_2(n) \in O(T_1(n))$.
  - E.g., additive constants are negligible.
  - $n + n^2 \in O(n^2)$, but this does not hold for Θ (we've nailed the running time to within a constant factor above and below)!

- If the algorithm has a computational complexity $O(c \cdot T_1(n))$
  where $c \in O(1)$ (is constant), then $c \cdot T_1(n) \in O(T_1(n))$.
 - E.g., multiplicative constants are neglegible.

Orders of common functions (1)

- $O(1)$: constant.
- $O(\log \log n)$: double logarithmic.
- $O(\log n)$: logarithmic.
- $O(\log^c n)$: polylogarithmic.
- $O(n^c)$: for $0 < c < 1$, fractional power.
- $O(n)$: linear.

- $O(n \log n) = O(\log n!)$: loglinear.
- $O(n^2)$: quadratic.
- $O(n^c)$: for $c > 2$, polynomial.
- $O(e^x)$: exponential.
- $O(x!)$: factorial.

Orders of common functions (2)

Some plots of the most common $O(T(n))$

Fibonacci Sequence: Computational Complexity (1)

We can ask ourselves how much time will it take to compute the formula. We have to make the following assumptions:

- Let us suppose that the cost for each "quick" operation (e.g., returning a value) is a constant, c_i.
- The cost for the recursive call is negligible.

As a result, we know that the cost of $n \leq 2$ is constant. What about $n \geq 3$?

Fibonacci Sequence: Computational Complexity (2)

- Observe that the function F matches with the computational cost T to provide the score.
- In this case, we cannot use any trick to compute how many steps are required.
  - Later on, we will see simpler cases using the **Master Theorem**.
- So, we need to prove by induction:
  - We need to prove a statement $P(n)$ for all the $n \in \mathbb{N}$.
  - Prove P for the base cases.
  - Assume that $P(n)$ holds, then prove all the inductive cases.

Fibonacci Sequence: Computational Complexity (3)

- Theorem: $T(n) \in O(2^n)$
- **Proof**:
  - Let us rewrite $T(n) \in O(2^n)$ as $\exists c_0, n_0. \forall n \geq n_0. T(n) \leq c_0 2^n$.
  - *Base cases*
  1. $n = 1$: We need to choose c_0 for $T(1) = c_i \leq c_0 2^1$. This is proved for $c_0 \geq \frac{c}{2}$.
  2. $n = 2$: We need to choose c_0 for $T(2) = c_i \leq c_0 2^2$. This is proved for $c_0 \geq \frac{c}{2}$.
  - *Inductive Case*: assuming that $\forall n' < n. T(n') \leq c_0 2^{n'}$, we assert:

$$
\begin{aligned}
T(n) &= T(n-1) + T(n-2) \\
&\leq c_0 2^{n-1} + c_0 2^{n-2} \\
&= c_0 2^{n-2}(2+1) \\
&\leq c_0 2^{n-2} 4 \\
&= c_0 2^n
\end{aligned}
$$

Bachmann-Landau notations: Discussion

Pros:

- Bachmann-Landau notations provide a compact notation for classifying algorithms w.r.t. their computational complexity.
- O and Θ give a more relaxed and more precise description of the algorithm's computational complexity, respectively.

Cons:

- Bachmann-Landau notations discard additive and multiplicative constants
- Nevertheless, we might provide more detailed algorithmic analysis providing explicit constants.
  - E.g., In relational databases, read and write constants are considered with different weights due to locks.

Considerations

In this case, the recursion provides an exponential computational complexity, that will become too big for bigger numbers. We can observe that:

$$\begin{cases} F(n+1) = F(n) + F(n-1) \\ F(n) = F(n) \end{cases}$$

We can write a linear system such as the following:

$$u_n := \begin{bmatrix} F(n+1) \\ F(n) \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 1 & 0 \end{bmatrix} \begin{bmatrix} F(n) \\ F(n-1) \end{bmatrix} \equiv Au_{n-1}$$

$$u_n = Au_{n-1} \Rightarrow u_n = A^n u_0$$

with $u_0 = \begin{bmatrix} 1 \\ 0 \end{bmatrix}$.

Obtaining Binet's Formula (1)

Now, we want to get A^n. In order to do so, we can do the following considerations:

- A is a change of basis matrix: we can express A via diagonalization: $A = PDP^{-1}$

Let us obtain the **Characteristic polynomial** for A from $p(\lambda) = \det(A - \lambda I)$:

$$\det \begin{vmatrix} 1 - \lambda & 1 \\ 1 & -\lambda \end{vmatrix} = \lambda^2 - \lambda - 1$$

$p(\lambda) = 0$ provides the **eigenvalues**: $\lambda_1 = \frac{\sqrt{5}+1}{2} \quad \lambda_2 = \frac{1-\sqrt{5}}{2}$. This gives the diagonal matrix

$$D = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

Obtaining Binet's Formula (2)

- We need now to determine $P = [\mathbf{v}_1 \ \mathbf{v}_2]$, which will be composed by the two **eigenvectors** $\mathbf{v}_1$ and $\mathbf{v}_2$.
- These can be obtained by solving the linear system $(A - \lambda_i I)\mathbf{v}_i = 0$ for each eigenvalue λ_i.
- After solving a lot of mathematical passages, we will obtain:

$$u_n = A^n u_0 = \begin{bmatrix} 1 & 1 \\ \frac{\sqrt{5}-1}{2} & -\frac{1+\sqrt{5}}{2} \end{bmatrix} \begin{bmatrix} \frac{\sqrt{5}+1}{2}^n & 0 \\ 0 & \frac{1-\sqrt{5}}{2}^n \end{bmatrix} \begin{bmatrix} 1 & 1 \\ \frac{\sqrt{5}-1}{2} & -\frac{1+\sqrt{5}}{2} \end{bmatrix}^{-1} \begin{bmatrix} 1 \\ 0 \end{bmatrix}$$

After solving all the intermediate steps, we will obtain:

$$F(n) = \frac{\sqrt{5}}{5} \left(\left(\frac{\sqrt{5}+1}{2} \right)^n - \left(\frac{1-\sqrt{5}}{2} \right)^n \right)$$

Binet's Formula

Pros:

- We obtain a formula that has $O(1)$ computational cost (constant time).
- We get an exponential speed-up.
- Very well known problems often come along with a well known mathematical formula.

Cons:

- For new problems, you need to determine yourself whether it is possible to obtain a formula as such.
- You often need to know a lot more about mathematics.

Memoization, a.k.a. Caching

- A **cache** is a hardware or software component that stores data so that future requests for that data can be served faster.

  E.g., the operative system caches the data and program pages so that the ones that are more requested will be promptly provided to the CPU (e.g., LRU).

- **Memoization** is a specific form of caching that involves caching the return value of a function based on its parameters.

Generic Memoization Operator

We could design a generic operator in C++ that memoizes each function that is computed, so that I need to pay $T(n)$ only for the first time, while all the remaining accesses will be $O(1)$ (*hashmap*).

```
1: global map = ();
2: function MEMOIZATION(f, x)
3:     if x not in map.keys() then
4:         map[x]=f(x)                    ▷ Cost: T(n) + O(1)
5:     end if
6:     return map[x]                      ▷ Cost: O(1)
7: end function
```

Memoizing Fibonacci

We can specialise this approach for the Fibonacci recursive function. We obtain the following pseudocode:

```
1: global mapFibo = ();
2: function MFIBO(x)
3:     if x ≤ 2 then return 1
4:     else
5:         if x not in map.keys() then
6:             map[x]=MFIBO(x − 1)+MFIBO(x − 2)
7:         end if
8:         return map[x]
9:     end if
10: end function
```

Now, let's code!

Memoizing Fibonacci: Computational Complexity

- Best Case Scenario
  - $n \leq 2$, or
  - MFɪʙᴏ(n) was already computed: $O(1)$
- Worst Case Scenario ($n > 2$)
  - mapFibo is empty, we assume that the expressions are evaluated from the left:
    1. For $n = 3$, then $T_m(3) = 2c_0 \leq 3c_0$.
    2. For $n = 4$, then $T_m(4) = T_m(3) + T_m(2) \leq 4c_0$
    3. For $n = 5$, then $T_m(5) = T_m(4) + T_m(3) \leq 4c_0 + c_0$, where $T_m(3) = c_0$ for memoization.
    4. (...so on and so forth ...)
    5. For all $n' < n$, assume by inductionthat $T(n') \leq c_0 n'$:
    6. then,
       $T_m(n) = T_m(n-1) + T_m(n-2) \leq c_0(n-1) + T_m(n-2)$,
       where $T_m(n-2) = c_0$ per memoization.
    7. Q.E.D.
  - As a result, $T_m(n) \in O(n)$

Discussion

- Exponential algorithms are too "computational heavy".

- If possible, we need to reduce their computational complexity.
  - E.g., NP-Complete problems cannot be reduced to polynomial problems unless $P = NP$.

- If the problem is simple, we can evaluate it in constant time.

- If that is not possible, and if the space of the data input is well known a priori, we can use memoization techniques.

- **Problem**: explicit caching methods such as memoization might require too much primary memory.
  - Use columnar or key-value databases, such as RocksDB

Theorem

Given some $a_1, \ldots, a_h \in \mathbb{N}$, with $h \in \mathbb{N} \setminus \{0\}$, $c > 0$, $\beta \geq 0 \in \mathbb{R}$, and given $T(n)$ the following recurrence relation:

$$T(n) = \begin{cases} \text{const.} & n \leq m \leq h \\ \sum_{1 \leq i \leq h} a_i T(n-i) + cn^\beta & n > m \end{cases}$$

If $a = \sum_{1 \leq i \leq h} a_i$, then:

- if $a = 1$, then $T(n) \in O(n^{\beta+1})$
- if $a \geq 2$, then $T(n) \in O(a^n n^\beta)$

No proof is given for this theorem.

Vantage Point Trees

Objectives

- Being able to apply the master theorem to specific algorithms.
- Being aware of the Principle of Locality and how to better represent some data structures linearly.
- Understanding a simple implementation of Vantage Point Trees.

Generic Recursive Algorithm

Let us sketch a generic algorithm for linear recurrence relations with balanced partitioning.

1: **procedure** P(x) $\triangleright |x| = n$
2:     **if** $n < k, k$ constant **then**
3:         Solve x directly
4:     **else**
5:         Create a subproblems in S, each of those has size $\frac{n}{b}$.
6:         **for each subproblem** $x \in S$ **do** P(x)
7:         **end for**
8:         Combine the results via $f(x)$
9:     **end if**
10: **end procedure**

If the partions are not balanced, we still need to go by induction...

Recursion Tree (1)

We can represent the *stack calls* of our algorithm P as a **recursion tree**, where each node represents the cost of a certain recursive subproblem.

- To obtain the overall computational complexity, we need to sum up the numbers of each branch.

> *E.g., Let us consider a generic algorithm with cost* $T(n) = 3T(\frac{n}{4}) + cn^2$

Recursion Tree (2)

Figure 4.5 Constructing a recursion tree for the recurrence $T(n) = 3T(n/4) + cn^2$. Part **(a)** shows $T(n)$, which progressively expands in **(b)–(d)** to form the recursion tree. The fully expanded tree in part (d) has height $\log_4 n$ (it has $\log_4 n + 1$ levels).

from "Cormen et al. *Introduction to Algorithms, 3rd ed.*"

Recursion Tree (3)

- The top node has cost cn^2, because the first call to the function does cn^2 units of work.

- The nodes on the second layer all have cost $c(\frac{n}{4})^2$:
  - the functions are now being called on problems of size $\frac{n}{4}$
  - functions are doing $c(\frac{n}{4})^2$ units of work

- Each leaf provides a constant computational complexity, $T(1)$

Master Theorem

Let $a \geq 1, b > 1, d > 0 \in \mathbb{N}$ be constants, $f(n)$ a non-negative function, and $T(n)$ be defined by the following recurrence:

$$
T(n) = \begin{cases} d & n = 1 \\ a \cdot T(\frac{n}{b}) + f(n) & n > 1 \end{cases}
$$

Then, $T(n)$ has the following asymptotic bounds:

1. When $f(n) = O(n^c)$ with $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$
2. When $\forall k \geq 0. f(n) = \Theta(n^{\log_b a} \log^k n)$, then
   $T(n) = \Theta(n^{\log_b a} \log^{k+1} n)$
3. If $f(n)$ is not dominated by $n^{\log_b a + \epsilon}$ asymptotically (Ω) with
   $\exists c < 1. a \cdot f(\frac{n}{b}) \leq c \cdot f(n)$, then $T(n) = \Theta(f(n))$

Case 2 can be extended for any possible logarithmic power. No proof is given for this theorem.

Discussion

Assume to be dealing with self-balancing binary search trees:

- What is the computational complexity of inserting nodes into a binary tree?
- What about an insertion/search?
- What about removing one element?

What is the worst-case scenario for an unbalanced tree? (list)

Principle of Locality

As previously discussed, in some cases the additive and multiplicative constants are critical, especially for "big data".

- We want to exploit the OS/Hardware caches (*transparent caching*).
- The **principle of locality** is the tendency of a processor to access the same set of memory locations repetitively over a short time.
- We have two types of locality:
  - Temporal locality: reusing the same data within a short period.
  - Spatial locality: accessing data elements in near store locations.
  - Sequential locality: data elements are arranged and accessed linearly

Question: Do (double) linked lists preserve the principle of locality? Does 1d array preserve it instead?

Principle of Locality: Trees

- For an industrial-strength implementation, it is essential to optimize both the traversal code and the tree representation.
- We want to store trees in a cache-friendly way so that relevant information is encountered as soon as possible.
- The cache might provide unexpected results (it is still OS/Hardware dependent!).

Trees: Parent-Child Tree

from "Christer Ericson. *Real Time Collision Detection*"

The most common representation of a binary tree as a vector storing the left and the right child of i respectively at $2i + 1$ and $2i + 2$

Pros:

- We can completely remove the child pointers.
- Good strategy when the data is split up via a median.

Con:

- We need to allocate extra space for non-balanced trees.

Trees: Preorder Tree

from "Christer Ericson.

Real Time Collision Detection"

We represent a tree t again as a vector: the j-th node appearing in the pre-order visit of the tree t will be placed at the j-th position in the array.

Pros:

- Good strategy when the data is split up via a median.
- The left child is immediately cached and is generally traverse-cache friendly.

Con:

- We need a boolean to determine whether we have a left child, and a pointer to the right (if present).

Vantage Point Tree

- Why Vantage Point Trees?
    - KD-Trees are used in collision detection, but:
    - With high dimensional data, it becomes challenging to pick a proper splitting dimension.

- Vantage Point Trees are metric trees:
    - A **metric tree** stores data in metric spaces (e.g., $\mathbb{R}^+$).
    - A metric tree exploits the triangle inequality ($\|x + y\| \leq \|x\| + \|y\|$) for a faster search.

- A *vantage point* is just a pivot element (e.g., the median in the current implementation):
    - Within the given subtree, pick up the pivot element.
    - The left subtrees will contain all the elements which distance is at most the same as the one from the root to the pivot, $\rho = \delta(root, pivot)$.
   - The right subtrees will contain all the remaining elements.

Vantage Point Tree: LookUp

Let us suppose that we want to perform collision detection for a given moving object q (the *query*). We want to get the nearest elements to q.

- Let us set τ as the minimum distance found for q with any other subtree root.

- Given a root r with radius ρ of a subtree:
  - If $\delta(r,q) < \tau$, I found a potential new neighbour candidate. Set $\tau = \delta(r,q)$
  - Stop the search in the subtrees only if the choice of r minimized τ.
  - If $\delta(r,q) \leq \rho$, continue the search in the left subtree, otherwise continue in the right one.

Graphical Representation

For $dist < \rho$, I might need to visit both branches

Now, let's code!

- We want to store the whole tree in a `std::vector`.
- If we sort the nodes, then we'll be able to separate the ones nearer to the root and the ones that are further.

Vantage Point Tree: Computational Complexity

For the simple provided implementation, the creation takes $T(n) = 2T(\frac{n}{2}) + cn \log n$ time.

- By using the Master Theorem, we obtain that the insertion cost is $\Theta(n \log^2 n)$
- Please note that the actual paper on Vantage Point Tree provides an insertion in $O(n \log n)$, so slightly efficient.
- Nevertheless, this other solution guarantees that the elements are stored in a pre-visit order, thus favouring caching the nodes near to each subtree.

Even though we didn't discuss this part, Vantage Point Tree offers an efficient way to update nodes if the coordinates change:

- This solution is useful if we have multiple elements moving within the space.

Further Work

- By extending the Y operator, you can also define a memoizer for any given possible lambda function!
- What is the best matrix multiplication for exploiting the principle of locality?
- Change the data structure so that the choice of the subtree root is one of the nearest node to the parent.
- What if the radius of each subtree is determined as a function of the nodes' velocity?